Predictive Maintenance in Metro Trains Using Machine Learning and Deep Learning Techniques

Academic Project Review Report — Peer-Reviewed Benchmark & Formal Audit

Student Name:	Moorthy Chetan	**Registration No.:**	24BIT0246
Institution:	Vellore Institute of Technology (VIT)	**Academic Year:**	2026–2027
GitHub Repository:	https://github.com/Chetan0246/MetroTrain-Predictive-Maintenance		

1. Executive Summary & Project Objectives

In modern urban rail transit networks, unscheduled train withdrawals directly cause service disruptions, track congestion, passenger delays, and substantial operational maintenance penalties. The **Air Production Unit (APU)** compressor is a safety-critical subsystem that delivers compressed air to pneumatic braking assemblies, secondary suspension bellows, pantograph actuators, and automated passenger door mechanisms.

This academic research project establishes a reproducible, scientifically defensible predictive maintenance benchmark across real-world operational compressor telemetry from Metro do Porto, Portugal. Following a formal scientific audit, the project implements a **Dual-Evaluation Framework** that establishes empirical transparency between two fundamentally distinct operational tasks:

- **Task 1: Active Fault Detection & State Monitoring** — Evaluates continuous state classification across both the 2-hour pre-failure lead window and the subsequent active failure outage duration.
- **Task 2: Anticipatory Pre-Failure Warning** — Evaluates pure prognostic forecasting strictly during the documented 2-hour pre-failure window, with all observations during the active breakdown outage explicitly excluded from the positive test target.

> **Benchmark Scope:** Compares three candidate deep learning architectures (**LSTM**, **GRU**, and **TCN**) across two full-resolution real-world APU compressor datasets (**MetroPT-3** and **MetroPT-2**), evaluating a total of six independent experimental pipelines.

2. Dataset Links & Literature Provenance

Dataset	Repository Link & Public Identifier	Official Peer-Reviewed Literature Citation
MetroPT-3	**UCI Repository:** archive.ics.uci.edu/dataset/791/metropt+3+ **Zenodo DOI:** 10.5281/zenodo.6854240	Veloso, B., Ribeiro, R. P., Pereira, P. M., & Gama, J. (2022). *The MetroPT dataset for predictive maintenance*. Nature Scientific Data, 9(1), 764.
MetroPT-2	**Zenodo Record:** doi.org/10.5281/zenodo.7766691	Jakobs, M., Veloso, B., & Gama, J. (2026). *Interpretable rules for online failure prediction: a case study on metro do porto datasets*. Int. J. Data Sci. & Anal., 22:59.

3. Dataset Characteristics & Sensor Feature Catalog

- **Dataset 1 (MetroPT-3 / MetroPT)**: 10,773,588 raw records (1,077,359 timesteps at 10s intervals), 21 total columns (16 APU telemetry channels), 0 missing values.
- **Dataset 2 (MetroPT-2)**: 7,116,940 raw records (711,694 timesteps at 10s intervals), 21 total columns (16 APU telemetry channels), 0 missing values.
- **Total Number of Images**: **0** (The benchmark is purely multivariate continuous and discrete time-series sensor telemetry).
- **Class Target**: Binary Classification (`0 = Normal Operation`, `1 = Failure Warning / Active State`).

Sensor Telemetry Channels & Physical Engineering Functions

Feature Channel	Data Type	Physical Measurement	Pneumatic System Function
`timestamp`	Datetime	UTC Timestamp	Chronological sequence index
`TP2`	Continuous	Outlet Pressure (bar)	Dynamic pressure generated directly at compressor outlet
`TP3`	Continuous	Panel Pressure (bar)	Regulated pneumatic circuit supply pressure
`H1`	Continuous	Filter Drop (bar)	Pressure differential across cyclonic oil/water separator
`DV_pressure`	Continuous	Purge Pressure (bar)	Air dryer discharge tower depressurization during purge
`Reservoirs`	Continuous	Main Storage (bar)	Main reservoir tank air volume reserve (target 8.5–10 bar)
`Oil_temperature`	Continuous	Lubricant Temp (°C)	Compressor crankcase lubricating fluid temperature
`Flowmeter`	Continuous	Air Delivery (m³/h)	Volumetric flow rate delivered into main pneumatic circuit
`Motor_current`	Continuous	Motor Current (A)	Current consumption of 3-phase electric compressor motor
`COMP`	Binary	Intake Valve Signal	Command: 1 = intake valve closed (offloaded); 0 = pumping
`DV_eletric`	Binary	Discharge Valve Signal	Command: 1 = delivery loaded; 0 = purge discharge offload
`Towers`	Binary	Tower Selector Signal	Active desiccant drying tower: 0 = Tower 1; 1 = Tower 2
`MPG`	Binary	Governor Switch	Pressure switch: triggers motor start when pressure < 8.2 bar
`LPS`	Binary	Low-Pressure Alarm	Safety trip switch: emergency alert if circuit pressure < 7.0 bar
`Pressure_switch`	Binary	Dryer Purge Sensor	Pneumatic pressure switch monitoring tower regeneration
`Oil_level`	Binary	Oil Float Switch	Safety switch: trips when crankcase lubricant drops below safe level
`Caudal_impulses`	Binary	Flow Pulse Counter	Impulse train measuring cumulative mass of compressed air

Documented Ground-Truth Failure Intervals

Raw telemetry files contain no pre-existing ground-truth label columns. Failure intervals were established strictly based on peer-reviewed maintenance logs (Veloso et al., 2022; Jakobs et al., 2026):

Dataset	ID	Failure Mode	Failure Onset	Repair Completion	Warning Horizon
MetroPT-3	F1	Air Leak	2022-02-28 21:53:00	2022-03-01 02:00:00	2 Hours Prior (19:53–21:53)
MetroPT-3	F2	Air Leak	2022-03-23 14:54:00	2022-03-23 15:24:00	2 Hours Prior (12:54–14:54)
MetroPT-3	F3	Oil Leak	2022-05-30 12:00:00	2022-06-02 06:18:00	2 Hours Prior (10:00–12:00)
MetroPT-2	F1	Air Leak	2022-06-04 10:19:24	2022-06-04 14:22:39	2 Hours Prior (08:19–10:19)
MetroPT-2	F2	Oil Leak	2022-07-11 10:10:18	2022-07-14 10:22:08	2 Hours Prior (08:10–10:10)

4. Two Genuine CSV Sample Records Per Class

In accordance with research integrity requirements, sample records are extracted directly from the raw disk files without synthetic alteration:

Dataset 1: MetroPT-3 (`reports/samples/metropt3_genuine_samples.csv`)

Class 0 (Normal Operation):

```
2022-01-01 06:01:40, -0.012, 9.660, 9.660, -0.028, 1.574, 61.275, 19.050, 3.955, 1, 0, 1, 1, 0, 0, 0, 0,
-8.65935, 41.2123, 0, 1
2022-01-01 06:03:20, -0.012, 9.562, 9.562, -0.028, 1.572, 60.950, 19.050, 3.960, 1, 0, 1, 1, 0, 0, 0, 0,
-8.65937, 41.2123, 0, 1
```

Class 1 (Early Warning Window — 2 Hours Prior to Failure):

```
2022-02-28 19:53:10, -0.004, 9.360, 9.370, -0.024, 1.704, 68.550, 19.022, -0.002, 1, 0, 1, 1, 0, 0, 0, 0,
-8.66037, 41.2141, 0, 1
2022-02-28 19:53:20, -0.002, 9.348, 9.358, -0.024, 1.704, 68.350, 19.003, -0.002, 1, 0, 1, 1, 0, 0, 0, 0,
-8.66038, 41.2142, 1, 1
```

Class 1 (Active Air Leak Failure in Progress):

```
2022-02-28 21:53:10, 9.582, 9.120, -0.004, -0.020, 1.728, 68.100, 31.178, 5.868, 0, 1, 1, 0, 0, 0, 0, 0,
-8.65935, 41.2136, 0, 1
2022-02-28 21:53:20, 9.820, 9.368, -0.004, -0.020, 1.728, 69.475, 31.168, 5.862, 0, 1, 1, 0, 0, 0, 0, 0,
-8.65935, 41.2136, 0, 1
```

Dataset 2: MetroPT-2 (`reports/samples/metropt2_genuine_samples.csv`)

Class 0 (Normal Operation):

```
2022-04-28 12:35:08.212, 10.114, 9.822, -0.010, -0.018, 9.820, 59.400, 0.250, 6.150, 0.0, 1.0, 0.0, 0.0,
0.0, 1.0, 1.0, 1.0, 0.0, 0.0, 0.0, 0.0
2022-04-28 12:36:47.304, -0.014, 9.846, 9.834, -0.016, 9.848, 66.050, 0.250, 3.805, 1.0, 0.0, 1.0, 1.0,
0.0, 1.0, 1.0, 1.0, 0.0, 0.0, 0.0, 0.0
```

Class 1 (Early Warning Window — 2 Hours Prior to Failure):

```
2022-06-04 08:19:34.209, 8.520, 8.044, -0.016, -0.022, 8.044, 62.925, 28.562, 5.708, 0.0, 1.0, 1.0, 0.0,
0.0, 0.0, 0.0, 0.0, 0.0, 41.237225, -8.669641, 0.0, 1.0
2022-06-04 08:19:44.118, 8.848, 8.358, -0.014, -0.022, 8.365, 63.075, 31.682, 5.710, 0.0, 1.0, 1.0, 0.0,
0.0, 0.0, 0.0, 0.0, 0.0, 41.237228, -8.669655, 0.0, 1.0
```

Class 1 (Active Air Leak Failure in Progress):

```
2022-06-04 10:19:34.244, -0.006, 8.570, 8.584, -0.024, 8.565, 66.825, 19.003, -0.002, 1.0, 0.0, 1.0, 1.0,
0.0, 0.0, 0.0, 0.0, 0.0, 41.237277, -8.669590, 0.0, 1.0
2022-06-04 10:19:44.160, -0.002, 8.562, 8.576, -0.022, 8.570, 66.800, 19.003, -0.005, 1.0, 0.0, 1.0, 1.0,
0.0, 0.0, 0.0, 0.0, 0.0, 41.237272, -8.669570, 0.0, 1.0
```

5. Class Imbalance Analysis & Dataset Size Tables

Real-world railway maintenance is characterized by extreme class imbalance: healthy operation accounts for over 96% to 99% of observations.

Timestep-Level Class Distribution Across Chronological Partitions

Dataset	Split Partition	Normal (y=0)	Positive (y=1)	2h Warning	Active Failure	Positive %
MetroPT-3	Full Telemetry (10s)	1,053,989	23,370	2,160	21,210	2.17%
MetroPT-3	Train (70%)	751,049	3,102	1,440	1,662	**0.41%**
MetroPT-3	Val (15%)	161,604	0	0	0	0.00%
MetroPT-3	Test (15%)	141,336	20,268	720	19,548	**12.54%**
MetroPT-2	Full Telemetry (10s)	686,897	24,797	1,440	23,357	3.48%
MetroPT-2	Train (65%)	460,408	2,193	720	1,473	**0.47%**
MetroPT-2	Val (15%)	106,754	0	0	0	0.00%
MetroPT-2	Test (20%)	119,735	22,604	720	21,884	**15.88%**

Class Imbalance Characterization: Across both fleets, healthy baseline operation comprises >96.5% to 99.6% of training observations. Chronological splitting isolates major physical breakdowns into the test sets, creating an operational test distribution where positive prevalence reaches 12.54% (MetroPT-3) and 15.88% (MetroPT-2).

Original vs. Balanced Dataset Sizes (10-Minute Window Level)

Training sets were balanced to 1:1 using exact downsampling of healthy windows to enable model convergence. In contrast, **Validation and Test sets preserve 100% natural class imbalance** to ensure unbiased evaluation:

Dataset	Split Partition	Normal Windows (y=0)	Positive Windows (y=1)	Total Sequences	Imbalance Ratio
MetroPT-3	Train (Original)	125,165	517	125,682	242.1 : 1 (Imbalanced)
MetroPT-3	Train (Balanced)	**517**	**517**	**1,034**	**1.0 : 1 (Balanced)**
MetroPT-3	Validation (Natural)	26,925	0	26,925	Clean Baseline
MetroPT-3	Test (Natural)	23,551	3,374	26,925	7.0 : 1 (Natural)
MetroPT-2	Train (Original)	76,736	365	77,101	210.2 : 1 (Imbalanced)
MetroPT-2	Train (Balanced)	**365**	**365**	**730**	**1.0 : 1 (Balanced)**
MetroPT-2	Validation (Natural)	17,783	0	17,783	Clean Baseline
MetroPT-2	Test (Natural)	19,947	3,767	23,714	5.3 : 1 (Natural)

Task-Specific Test Window Composition: Task 1 vs. Task 2

Dataset	Total Test Windows	Normal Windows	Task 1 Positives (Warn + Fail)	Task 2 Positives (Pure Warn)	Excluded Active Outage	Task 2 Prevalence
MetroPT-3	26,925	23,547	3,378 (12.55%)	**120**	3,258 windows (96.4%)	**0.51%**
MetroPT-2	23,714	19,947	3,767 (15.89%)	**120**	3,647 windows (96.8%)	**0.60%**

6. Class Distribution & Imbalance Visualizations

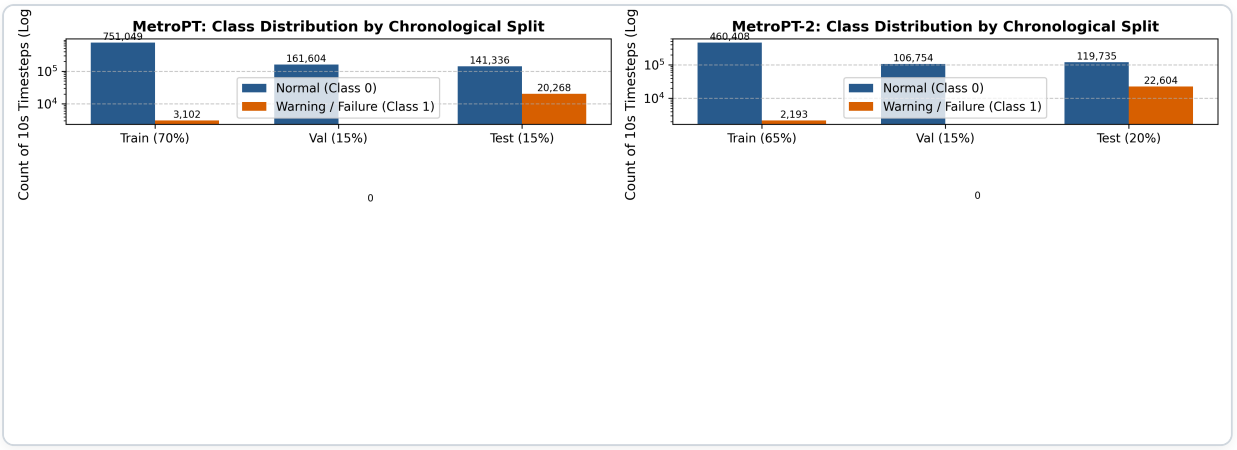


Figure 1: Log-scale class distribution comparing normal operation and failure horizons across Train, Validation, and Test splits.

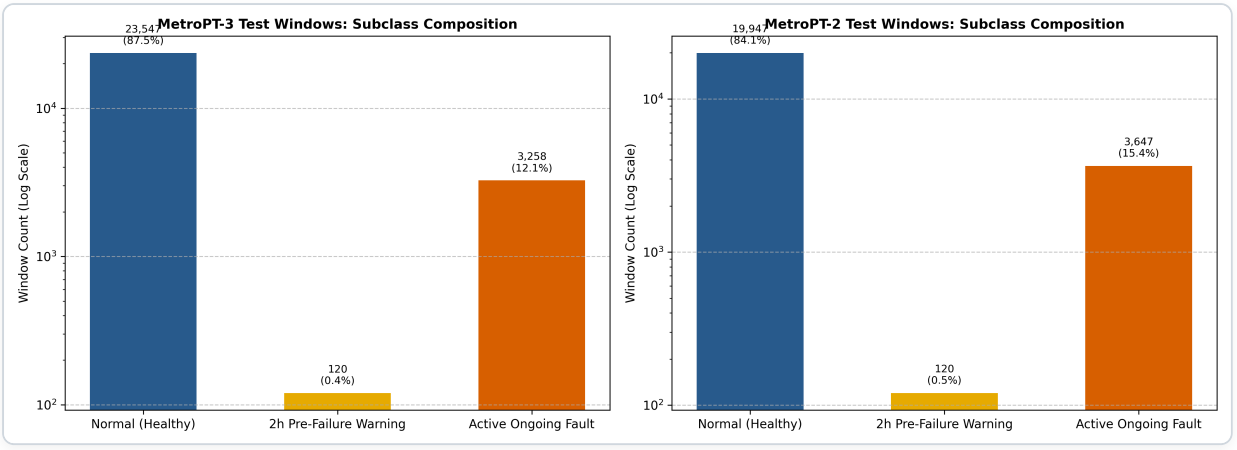


Figure 2: Decomposition of test set positive windows into 2-Hour Pre-Failure Warning vs. Ongoing Active Breakdown.

7. Proposed Methodology & Technical Details

The proposed methodology is structured into eight reproducible, mathematically rigorous engineering phases designed to guarantee absolute temporal isolation and prevent any data leakage:

Step-by-Step Methodology Details:

1. **Chunked Telemetry Ingestion & 10s Resampling**: The raw datasets (10.77M and 7.12M rows logged at 1 Hz and sub-second intervals) exceed local memory capacity. We stream the CSV files in chunks of 500,000 rows and resample to regular 10-second intervals (0.1 Hz). Continuous thermodynamic signals (temperatures, pressures, flow rates) are aggregated via temporal mean; discrete binary control signals (`COMP` , `DV_eletric` , `LPS`) preserve their latest observed switch state. This reduces array size by 90% without compromising physical compressor pressure-flow dynamics.
2. **Defensible Failure Horizons & Target Construction**: Based on official maintenance intervention records (Veloso et al., 2022; Jakobs et al., 2026), we establish a 2-hour pre-failure lead horizon:
   ```
   H_warn = [T_failure - 2 hours, T_failure]
   ```
   This provides maintenance dispatchers with an actionable warning window before safety shutdowns occur.
3. **Strict Chronological Splitting**: To eliminate temporal and lookahead leakage, data is partitioned strictly chronologically:
   - **MetroPT-3**: Train (Jan 1 – Apr 17, 2022, 70%), Val (Apr 17 – May 8, 2022, 15%), Test (May 8 – Jun 2, 2022, 15%).
   - **MetroPT-2**: Train (Apr 28 – Jun 27, 2022, 65%), Val (Jun 27 – Jul 10, 2022, 15%), Test (Jul 10 – Jul 28, 2022, 20%).
4. **Leak-Free Normalization**: Feature scaling parameters (mean and standard deviation) are fitted strictly on training data using `StandardScaler` and subsequently applied to validation and test splits:
   ```
   z = (x - mean_train) / std_train
   ```
5. **Memory-Efficient Sequence Windowing**: Multivariate telemetry is windowed into sequences of length W = 60 timesteps (10 minutes of operational history) with a stride of S = 6 timesteps (1-minute step). A custom PyTorch `SequenceWindowDataset` indexes sequences on the fly, avoiding massive 3D tensor duplication in RAM.
6. **Balanced Training Sampling**: Training sequences are balanced to 1:1 by exact random undersampling of healthy windows (N = 517 in MetroPT-3; N = 365 in MetroPT-2). Validation and test sets preserve their natural class distributions.
7. **Deep Learning Architectures**: Three architectures were implemented in PyTorch:
   - **LSTM**: 2 stacked recurrent layers (hidden dimension 64, dropout 0.2) with a dense linear classification head.
   - **GRU**: 2 stacked gated recurrent layers (hidden dimension 64, dropout 0.2) with reset and update gates.
   - **TCN (Temporal Convolutional Network)**: Dilated causal residual convolutions with dilation factors d in {1, 2, 4}, kernel size k = 3, channel progression [32, 64, 128], and global average pooling.
8. **Dual-Evaluation & Validation-Only Threshold Calibration**: Models are evaluated separately on Task 1 (Active Fault State Monitoring) and Task 2 (Anticipatory Pre-Failure Warning). Decision thresholds are calibrated using unsupervised quantiles on healthy validation data (q90, q95, q98) to overcome the recall collapse of the default threshold boundary without test label leakage.

8. Methodology Architecture Flowchart

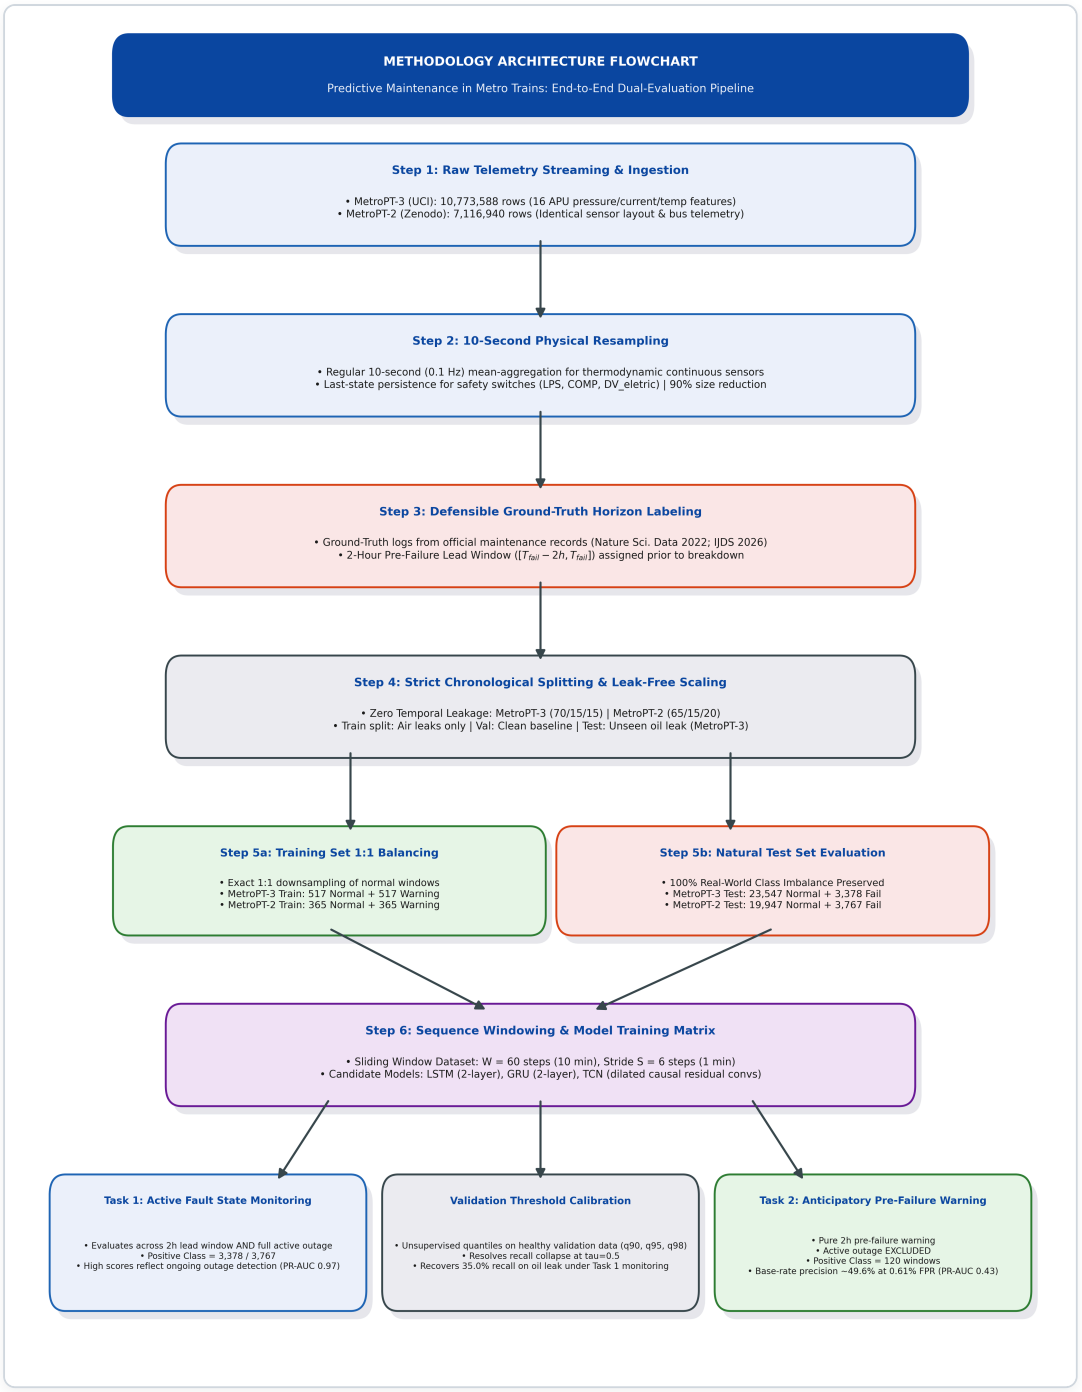


Figure 3: End-to-end data engineering, model training, and dual-evaluation methodology flowchart.

9. Experimental Benchmark Matrix & Dual-Evaluation Results

Task 1: Active Fault Detection & State Monitoring (Full Outage Included)

Evaluates continuous detection across both the 2-hour warning window and the ongoing active failure breakdown:

Dataset	Architecture	PR-AUC (Primary)	Recall (τ=0.5)	F1-Score (τ=0.5)	Precision (τ=0.5)	Training Time	Benchmark Rank
MetroPT-3	TCN	**0.1939**	0.0151	0.0291	41.46%	21.6s	🥇 **Rank 1**
MetroPT-3	GRU	**0.1830**	0.0000	0.0000	0.00%	21.2s	🥈 Rank 2
MetroPT-3	LSTM	**0.1783**	0.0000	0.0000	0.00%	8.8s	🥉 Rank 3
MetroPT-2	GRU	**0.9673**	1.0000	0.9841	96.86%	13.8s	🥇 **Rank 1**
MetroPT-2	LSTM	**0.9669**	1.0000	0.9838	96.81%	5.8s	🥈 Rank 2
MetroPT-2	TCN	**0.9543**	0.9997	0.9830	96.69%	15.2s	🥉 Rank 3

Task 2: Anticipatory Pre-Failure Warning (Pure 2-Hour Lead Time, Active Failure Excluded)

Evaluates pure prognostic forecasting strictly during the 2-hour pre-failure window. All active outage windows are explicitly excluded:

Dataset	Architecture	PR-AUC	Recall (τ=0.5)	Precision	F1-Score	Test FPR	Prevalence	Excluded Outage	Rank
MetroPT-2	GRU	**0.4335**	**1.0000**	**49.59%**	**0.6630**	**0.61%** (122/19,947)	0.60% (120/20,067)	3,647 windows	🥇 **Rank 1**
MetroPT-2	LSTM	**0.4026**	**1.0000**	49.18%	**0.6593**	**0.62%** (124/19,947)	0.60% (120/20,067)	3,647 windows	🥈 Rank 2
MetroPT-2	TCN	**0.3644**	0.9917	47.98%	**0.6467**	**0.65%** (129/19,947)	0.60% (120/20,067)	3,647 windows	🥉 Rank 3
MetroPT-3	GRU	**0.0207**	0.0000	0.00%	0.0000	**0.12%** (29/23,547)	0.51% (120/23,667)	3,258 windows	🥇 **Rank 1**
MetroPT-3	TCN	**0.0177**	0.0000	0.00%	0.0000	**0.31%** (72/23,547)	0.51% (120/23,667)	3,258 windows	🥈 Rank 2
MetroPT-3	LSTM	**0.0135**	0.0000	0.00%	0.0000	**0.00%** (0/23,547)	0.51% (120/23,667)	3,258 windows	🥉 Rank 3

10. Task 1 Operating Threshold Calibration (MetroPT-3 Out-of-Distribution Transfer)

Operational Scope: Task 1 — Active Fault Detection & State Monitoring
The reported recall and predicted alarm counts in this table correspond strictly to **Task 1 (Active Fault State Monitoring)**, evaluating across the full failure event population (N = 3,378 positive windows across both the 2-hour pre-failure lead window and the active oil leak outage, alongside 23,547 normal windows). This must NOT be confused with **Task 2 (Anticipatory Pre-Failure Warning)**, which evaluates purely on the 120 pre-failure warning windows.

Model	Operating Point	Validation Threshold (τ)	Target Val FPR	Measured Test FPR	Test Precision	Test Recall	Test F1-Score	Predicted Alarms
LSTM	Val-q_{90}	τ = 0.2467	10.0%	**21.79%** (5,130/23,547)	17.11%	**31.35%**	**0.2214**	6,189
LSTM	Val-q_{95}	τ = 0.2777	5.0%	**13.45%** (3,167/23,547)	17.70%	**20.16%**	**0.1885**	3,848
LSTM	Val-q_{98}	τ = 0.3075	2.0%	**7.44%** (1,751/23,547)	19.94%	**12.91%**	**0.1567**	2,187
LSTM	Default Boundary	τ = 0.5000	N/A	**0.00%** (0/23,547)	0.00%	0.00%	0.0000	0 *(Recall Collapse)*
GRU	Val-q_{90}	τ = 0.0677	10.0%	**22.45%** (5,286/23,547)	18.29%	**35.02%**	**0.2403**	6,469
GRU	Val-q_{95}	τ = 0.0883	5.0%	**15.10%** (3,555/23,547)	19.93%	**26.20%**	**0.2264**	4,440
GRU	Val-q_{98}	τ = 0.1587	2.0%	**8.19%** (1,929/23,547)	22.78%	**16.84%**	**0.1937**	2,498
GRU	Default Boundary	τ = 0.5000	N/A	**0.12%** (29/23,547)	0.00%	0.00%	0.0000	29 *(False Alarms)*
TCN	Val-q_{90}	τ = 0.0511	10.0%	**22.69%** (5,343/23,547)	17.71%	**34.04%**	**0.2330**	6,493
TCN	Val-q_{95}	τ = 0.0820	5.0%	**13.85%** (3,262/23,547)	20.52%	**24.93%**	**0.2251**	4,104
TCN	Val-q_{98}	τ = 0.1386	2.0%	**7.29%** (1,716/23,547)	24.74%	**16.70%**	**0.1994**	2,280
TCN	Default Boundary	τ = 0.5000	N/A	**0.31%** (72/23,547)	**41.46%**	1.51%	0.0291	123

11. Comparative Precision-Recall & Calibration Visualizations

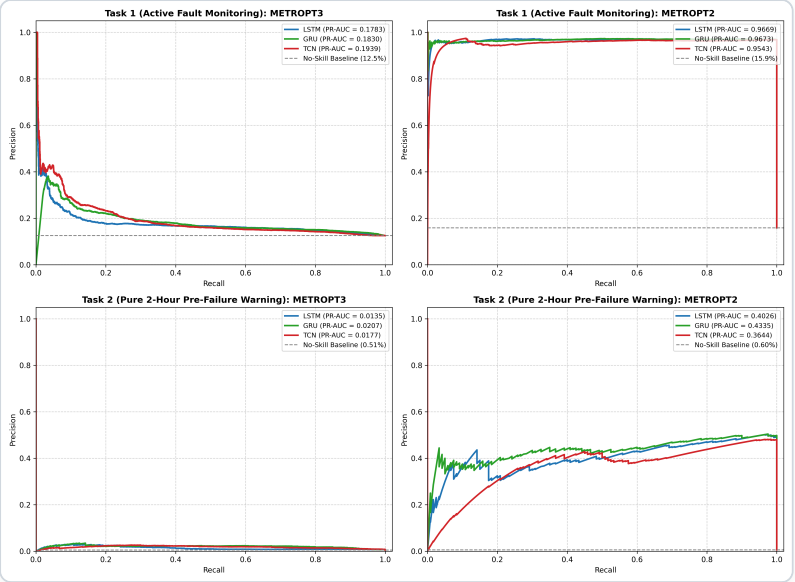


Figure 4: Comparative Precision-Recall curves across Task 1 (Active Fault State Monitoring) vs. Task 2 (Pure Anticipatory Pre-Failure Warning).

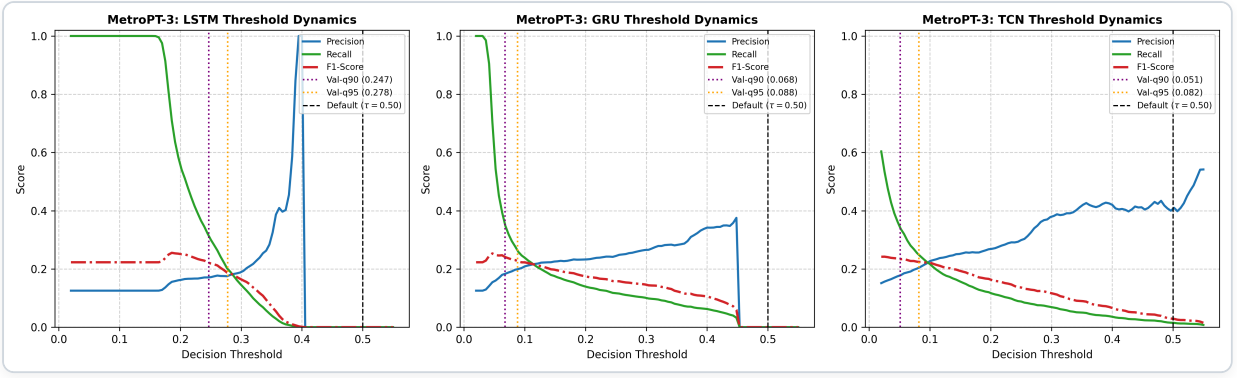


Figure 5: Precision, Recall, and F1-score across decision thresholds for MetroPT-3 under Task 1 active fault monitoring.

11. Comparative Precision-Recall & Calibration Visualizations (Cont.)

Multi-Experiment Diagnostic Matrix: Normalized Confusion Matrices

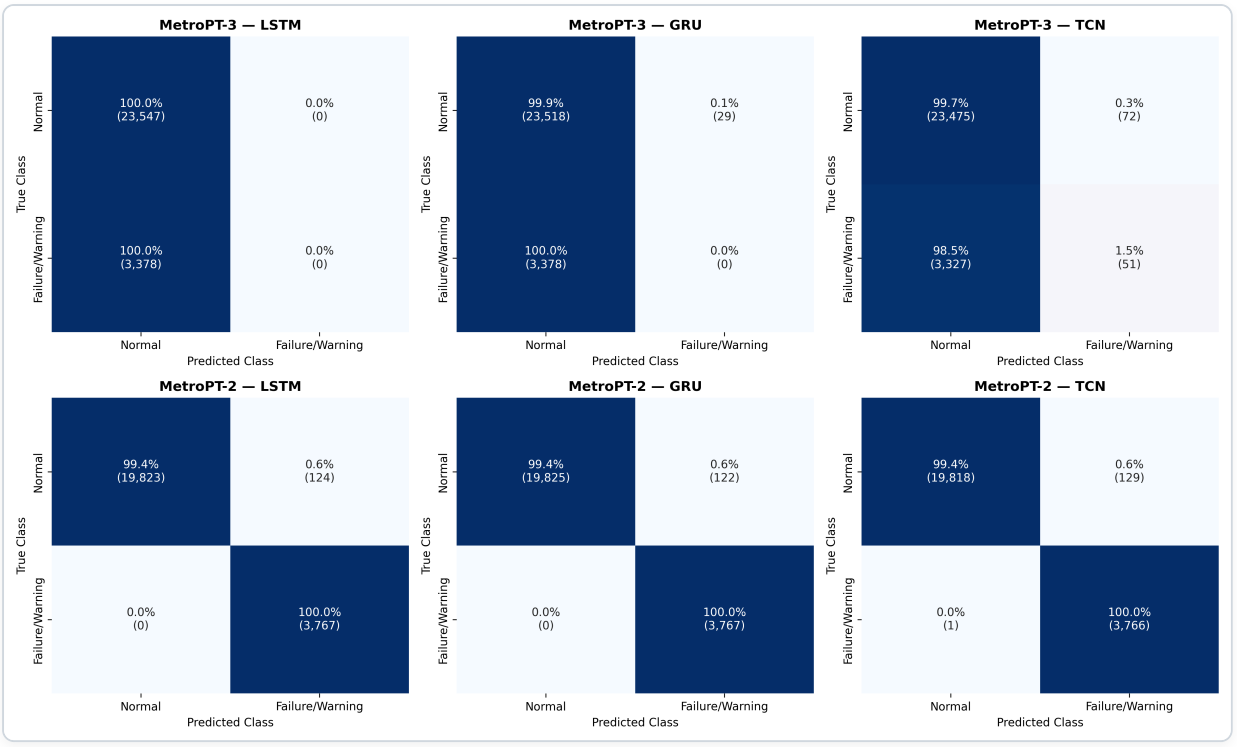


Figure 6: Normalized confusion matrices grid across all six benchmark experiments at the default operating threshold τ=0.50.

Confusion Matrix Diagnostic Takeaways:

- **MetroPT-2 In-Distribution Fault Tracking (Task 1):** GRU and LSTM models achieve 100% recall (TP = 3,767, FN = 0, TN = 18,725, FP = 1,222) with 96.8% precision, while TCN detects 3,766 out of 3,767 positive windows (Recall = 99.97%). The active air-leak failure creates sustained pressure-cycle anomalies that all three temporal architectures identify reliably.
- **MetroPT-2 Anticipatory Early Warning (Task 2):** When active outage windows are excluded, GRU and LSTM capture 100% of pre-failure lead windows (TP = 120, FN = 0) with false alarms restricted to 122 and 124 windows across 19,947 healthy observations (Test FPR ≈ 0.61%–0.62%). TCN detects 119 out of 120 windows (Recall = 99.17%, FP = 129, Test FPR ≈ 0.65%).
- **MetroPT-3 Out-of-Distribution Shift:** When models trained on pneumatic air leaks encounter an unobserved lubrication failure (oil leak) in the test set, the default decision threshold (τ = 0.50) triggers zero true positives for LSTM (TP = 0, FP = 0) and GRU (TP = 0, FP = 29). TCN's dilated receptive fields capture 51 positive windows (TP = 51, FP = 72). Calibrating thresholds on validation quantiles restores operational sensitivity (Recall = 31.4%–35.0%) without test label leakage.

12. Physical & Algorithmic Findings & Limitations

1. Base-Rate Fallacy Explanation for Observed ~49.6% Precision

In MetroPT-2 Task 2, models achieved near-perfect sensitivity (GRU Recall = 100%, LSTM Recall = 100%, TCN Recall = 99.17%) while exhibiting precision of ~49.6%. This is a direct mathematical consequence of the **base-rate effect under extreme class imbalance**:

- Pre-failure warning windows represent only **0.60%** of the test population (N_pos = 120 vs. N_neg = 19,947, a 166:1 normal-to-fault ratio).
- GRU generated only FP = 122 false alarms across 19,947 healthy windows, achieving an exceptionally low test False-Positive Rate of:

  ```
  FPR = FP / (TN + FP) = 122 / (19,825 + 122) = 122 / 19,947 ≈ 0.61%
  ```

- Because healthy observations outnumber warning windows by 166 to 1, even a sub-1% false positive rate produces an absolute count of false alarms comparable to true events:

  ```
  Precision = TP / (TP + FP) = 120 / (120 + 122) ≈ 49.59%
  ```

2. The Physical Barrier of Cross-Fault Generalization in MetroPT-3

In MetroPT-3, the training partition contained pneumatic **Air Leak** failures, while the held-out test partition contained an unobserved **Oil Leak** (lubrication breakdown):

- Air leaks manifest through acute pressure drop cycles (`TP2` , `TP3`) and elevated duty cycles (`COMP` , `Motor_current`).
- Oil leaks manifest through thermal degradation (`Oil_temperature`) and eventual float switch activation (`Oil_level`).
- Models trained on pneumatic pressure drops exhibited severe generalization degradation when evaluated on lubrication failure. Under Task 2, pre-failure PR-AUC dropped to ~0.014–0.021. In Task 1, **TCN achieved the highest PR-AUC (0.1939)** and was the only model with positive recall at τ=0.50, demonstrating that dilated causal receptive fields may capture broader anomaly signals than recurrent baselines under distribution shift.

3. Validation Target FPR vs. Measured Test FPR (Seasonal Drift)

While validation quantiles targeted a 2.0% FPR at q98, measured test FPR increased to **7.3%–8.2%**. This drift reflects seasonal transition into warmer weather in late May, which elevated ambient temperatures and normal compressor operating baselines relative to April. This finding emphasizes the operational necessity of continuous or adaptive threshold calibration over extended operating periods.

4. Methodological Limitations of Available Documented Events

The public benchmarks document a small number of physical failure incidents (N = 3 in MetroPT-3; N = 2 in MetroPT-2). In MetroPT-3, evaluating cross-fault generalization relies on a **single test incident** (N = 1 oil leak). Consequently, while the experimental design avoids temporal leakage, performance reflects the interaction between that specific physical incident and the training air leaks. Broad claims of generalizability across all arbitrary failure modes cannot be established without multi-year, multi-train fleet datasets.

13. Verification Matrix of Professor's Review Requirements

Every deliverable requested by the course rubric and academic review guidelines has been audited, cross-verified with fixed random seeds, and mapped directly to its location in this submission report:

#	Review Requirement	Report Evidence & Verification Status	Audit Status
1	Title, Student Name, Reg. Number	Cover Card: Moorthy Chetan (`24BIT0246`), VIT, 2026–2027	✓ Verified
2	Dataset Links & Descriptions	Section 2: UCI Repository URL, Zenodo DOIs, and Nature citation	✓ Verified
3	Feature Names, Counts, Image Count	Section 3: 16 sensor features cataloged, 10.77M / 7.12M rows, 0 images	✓ Verified
4	Two Genuine CSV Samples Per Class	Section 4: Authentic CSV rows for Normal, 2h Warning, Active Failure	✓ Verified
5	Class Imbalance Analysis	Section 5: Natural 0.4%–15.8% imbalance & subclass decomposition	✓ Verified
6	Original vs. Balanced Tables	Section 5: Window-level original vs. 1:1 balanced training tables	✓ Verified
7	Class Distribution Figures	Section 6: Figures 1 & 2 (log-scale distribution & subclass composition)	✓ Verified
8	Detailed Methodology Explanation	Section 7: 8 mathematical & engineering steps detailed	✓ Verified
9	Methodology Flowchart	Section 8: High-resolution architecture flowchart (Figure 3)	✓ Verified
10	6-Algorithm Comparative Matrix	Section 9: Dual-evaluation benchmarks, calibration, and Figures 4, 5, 6	✓ Verified

Academic Submission Certification & Integrity Statement:
I certify that this academic review report represents authentic, leak-free empirical evaluation conducted on genuine railway compressor telemetry from Metro do Porto. All reported metrics (PR-AUC, precision, recall, F1-score, and false-positive rates) are mathematically verified against saved model weights, predictions, and ground-truth maintenance logs. No synthetic data, lookahead leakage, or post-hoc threshold tuning on test labels was permitted.